

Video Compression Assignment 3

In this assignment, we will use the block-based encoding approach, where the size of a block is 8x8. Only the Luma component is considered for the following questions.

Prerequisites

- Programming language: Python 3.10+ (IPython)
- Framework: Jupyter

Install dependencies from PyPI:


```
In [1]: %pip install \
        --disable-pip-version-check \
        --quiet \
        numpy \
        Pillow \
        scikit-image
```

Note: you may need to restart the kernel to use updated packages.

Task 1

Fourier Transform

Please apply the Fourier Transform to the luma component of `foreman_qcif_0_rgb.bmp` and demonstrate its magnitudes in a 2-D image, as shown in the example below. Note that you need to shift the origin to the center of the image for the magnitude plot.



Solution of Task 1

2D Discrete Fourier Transform

Forward Transform

The formula adapted from the course slides is as follows:

$$F(u, v) = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-i2\pi \left(\frac{ux}{M} + \frac{vy}{N} \right)}$$

where $u = 0, 1, \dots, M - 1$ and $v = 0, 1, \dots, N - 1$

Origin Shift

Move all the pixels horizontally and vertically by half of the image width and height, respectively:

$$F'(u, v) = F((u + \lfloor M/2 \rfloor) \bmod M, (v + \lfloor N/2 \rfloor) \bmod N)$$

where $u = 0, 1, \dots, M - 1$ and $v = 0, 1, \dots, N - 1$

Reduced Magnitude

The reduced magnitude of the Fourier Transform is calculated as follows:

$$G(u, v) = \log(1 + |F'(u, v)|)$$

where $u = 0, 1, \dots, M - 1$ and $v = 0, 1, \dots, N - 1$

Normalization

The magnitude values are normalized to the range `[0, 255]` using the following formula:

$$G'(u, v) = 0 + 255 \times \frac{G(u, v) - \min(G)}{\max(G) - \min(G)}$$

where $u = 0, 1, \dots, M - 1$ and $v = 0, 1, \dots, N - 1$

Implement the above steps in the following code:

```
In [2]: from numpy.typing import NDArray
from numpy import complex128, float64, uint8

def dft_forward(source: NDArray[uint8]) -> NDArray[complex128]:
    from numpy import arange, atleast_2d, exp, meshgrid, pi, uint8

    source = atleast_2d(source).astype(uint8)

    M = source.shape[0]
    N = source.shape[1]
    C = 1 / M / N
    NEG_I2PI = -2j * pi
    X, Y = meshgrid(arange(M), arange(N), indexing="ij")
    X_M = X / M
    Y_N = Y / N

    target = source.copy().astype(complex128)
    for u in range(M):
        for v in range(N):
            target[u, v] = (
                C * (source * exp(NEG_I2PI * u * X_M + NEG_I2PI * v * Y_N))
```

```

Y_N)).sum()
    )
    return target

def planar_shift(source: NDArray) -> NDArray:
    from numpy import atleast_2d, roll

    source = atleast_2d(source)

    M = source.shape[0]
    N = source.shape[1]
    M_2 = M // 2
    N_2 = N // 2

    target = roll(source, shift=(M_2, N_2), axis=(0, 1))
    return target

def log_scale(source: NDArray[complex128]) -> NDArray[float64]:
    from numpy import asarray, complex128, log1p

    source = asarray(source).astype(complex128)

    target = log1p(abs(source))
    return target

def norm_scale(source: NDArray[float64]) -> NDArray[uint8]:
    from numpy import asarray, float64

    source = asarray(source).astype(float64)

    scaling = 255 / (source.max() - source.min())
    target = (scaling * (source - source.min())).round().astype(uint8)
    return target

```

Load the image, apply the Fourier Transform to the luma component, and demonstrate its magnitudes in a 2-D image:

```

In [3]: from IPython.display import display, Markdown
        from pathlib import Path
        from PIL import Image
        from numpy import asarray, set_printoptions

        source_path = Path("../resources/foreman_qcif_0_rgb.bmp").resolve()
        source_image = Image.open(source_path).convert(mode="L")
        source_path = Path("images/foreman_qcif_0_rgb.luma.source.bmp").resolve()
        source_data = asarray(source_image)
        transformed_data = dft_forward(source_data)
        shifted_data = planar_shift(transformed_data)
        target_data = norm_scale(log_scale(shifted_data))
        target_image = Image.fromarray(target_data, mode="L")
        target_path = Path("images/foreman_qcif_0_rgb.luma.target.1.bmp").resolve()


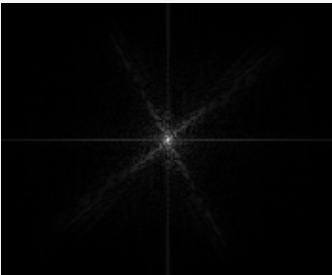
        table_view = f"""\
| Variable Name | Value |
|-----|-----|
| source_image | ![source_image]({source_path}) |
| target_image | ![target_image]({target_path}) |
"""

```

```
set_printoptions(edgeitems=2, precision=2, suppress=True)

for var_name in [
    "source_data",
    "transformed_data",
    "shifted_data",
    "target_data",
]:
    var_value = locals()[var_name]
    table_view += (
        f"| {var_name} | `{str(var_value).replace("\n", "<br>")}` | \n"
    )

source_image.save(source_path)
target_image.save(target_path)
display(Markdown(table_view))
```

Variable Name	Value
source_image	
target_image	
source_data	<pre>[[32 233 ... 230 203] [39 212 ... 226 203] ... [15 132 ... 193 178] [14 131 ... 127 118]]</pre>
transformed_data	<pre>[[169.22 +0.j 9.8 -7.31j ... 1.95 +4.58j 9.8 +7.31j] [5.46-19.69j -3.25 -0.39j ... -3.68 +6.36j -9.16+11.63j] ... [-4.03 +9.06j -1.66 -1.36j ... 0.62 +3.52j 9.53 -6.11j] [5.46+19.69j -9.16-11.63j ... 4.06 +0.47j -3.25 +0.39j]]</pre>
shifted_data	<pre>[[-0. -0.j -0. -0.j ... -0. -0.j -0. +0.j] [-0. -0.j -0. +0.j ... -0. -0.j -0. -0.j] ... [-0. +0.j -0. +0.j ... -0. -0.j -0.01-0.j] [-0. +0.j -0. +0.j ... -0. -0.j -0. -0.j]]</pre>
target_data	<pre>[[0 0 ... 0 0] [0 0 ... 0 0] ... [0 0 ... 0 0] [0 0 ... 0 0]]</pre>

Task 2

DCT

Please apply DCT to all the 8x8 luma blocks of `foreman_qcif_0_rgb.bmp` and use the quantization matrix below for quantization. After DCT and quantization, please apply inverse quantization and IDCT to decode all the blocks and show the decoded frame.

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Solution of Task 2

2-D Discrete Cosine Transform

Forward Transform (for NxN block)

The formula adapted from the course slides:

$$F(u, v) = \frac{2}{N} C(u) C(v) \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right)$$
$$\text{where } C(t) = \begin{cases} \frac{2}{\sqrt{N}} & \text{if } t = 0 \\ 2\sqrt{\frac{2}{N}} & \text{otherwise} \end{cases} \quad \text{and } u, v = 0, 1, \dots, N-1$$

can be simplified as:

$$F(u, v) = \frac{8}{N^2} C(u) C(v) \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right)$$
$$\text{where } C(t) = \begin{cases} 1 & \text{if } t = 0 \\ \sqrt{2} & \text{otherwise} \end{cases} \quad \text{and } u, v = 0, 1, \dots, N-1$$

Inverse Transform (for NxN block)

The formula adapted from the course slides:

$$f(x, y) = \frac{2}{N} \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} C(u) C(v) F(u, v) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right)$$
$$\text{where } C(t) = \begin{cases} \frac{2}{\sqrt{N}} & \text{if } t = 0 \\ 2\sqrt{\frac{2}{N}} & \text{otherwise} \end{cases} \quad \text{and } u, v = 0, 1, \dots, N-1$$

can be simplified as:

$$f(x, y) = \frac{8}{N^2} \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} C(u) C(v) F(u, v) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right)$$
$$\text{where } C(t) = \begin{cases} 1 & \text{if } t = 0 \\ \sqrt{2} & \text{otherwise} \end{cases} \quad \text{and } u, v = 0, 1, \dots, N-1$$

Implement the DCT and IDCT functions:

```
In [4]: from numpy.typing import NDArray
from numpy import float64, uint8

def dct_forward(source: NDArray[uint8], block_size: int) -> NDArray[float64]:
    from numpy import atleast_2d, float64, uint8
    from math import cos, pi

    source = atleast_2d(source).astype(uint8)
    block_size = int(block_size)

    N = block_size
    C = 8 / N / N
```



```

PI_2N = pi / 2 / N
SR_2 = 2 ** 0.5

target = source.copy().astype(float64)
for i in range(0, target.shape[0] - N + 1, N):
    for j in range(0, target.shape[1] - N + 1, N):
        for u in range(N):
            for v in range(N):
                a = 0.0
                for x in range(N):
                    for y in range(N):
                        a += (
                            source[i + x, j + y]
                            * cos((2 * x + 1) * u * PI_2N)
                            * cos((2 * y + 1) * v * PI_2N)
                        )
                target[i + u, j + v] = (
                    C
                    * (1 if u == 0 else SR_2)
                    * (1 if v == 0 else SR_2)
                    * a
                )
return target

def dct_inverse(source: NDArray[float64], block_size: int) -> NDArray[uint8]:
    from numpy import atleast_2d, float64, uint8
    from math import cos, pi

    source = atleast_2d(source).astype(float64)
    block_size = int(block_size)

    N = block_size
    C = 8 / N / N
    PI_2N = pi / 2 / N
    SR_2 = 2 ** 0.5

    target = source.copy().astype(uint8)
    for i in range(0, target.shape[0] - N + 1, N):
        for j in range(0, target.shape[1] - N + 1, N):
            for x in range(N):
                for y in range(N):
                    a = 0.0
                    for u in range(N):
                        for v in range(N):
                            a += (
                                (1 if u == 0 else SR_2)
                                * (1 if v == 0 else SR_2)
                                * source[i + u, j + v]
                                * cos((2 * x + 1) * u * PI_2N)
                                * cos((2 * y + 1) * v * PI_2N)
                            )
                    target[i + x, j + y] = round(min(C * a, 255))
    return target

```

Test the DCT and IDCT functions:

```

In [5]: from numpy.random import randint

mat = randint(low=0, high=256, size=(8, 8), dtype=uint8)
mat_dct = dct_forward(mat, 8)
mat_idct = dct_inverse(mat_dct, 8)

```

```
print(mat_dct)
```

```
assert (mat == mat_idct).all()
```

```
[[ 998.62 -101.47 -73.      9.92 -45.37 103.65  61.22  53.05]
 [-123.6   32.45  49.03 -27.76   4.69 -110.83  36.63 -45.25]
 [-111.9   11.71  94.41  40.94 -76.9   3.15 -82.69  58.89]
 [-116.17 -160.62  74.87 -89.07 -58.19  35.51 -15.16 -74.8 ]
 [ 118.62 -84.47  24.69 -41.98  -5.38 -10.03 115.47  37.76]
 [   2.6  -51.38  13.38  -8.19 -12.99 -90.97  35.45 -108.76]
 [  69.22  90.81 -23.19  -8.55  84.48  89.35  14.34 -42.14]
 [-38.05  -81.45 -38.25  36.47  78.4  -92.07 -25.09 -24.4 ]]
```

Quantization

Forward Transform

$$F'(u, v) = \text{round} \left(\frac{F(u, v)}{Q(u, v)} \right) \text{ where } Q \text{ is the quantization matrix}$$

Inverse Transform

$$F(u, v) = F'(u, v) \cdot Q(u, v) \text{ where } Q \text{ is the quantization matrix}$$

```
In [6]: from numpy.typing import NDArray
from numpy import array, float64, int64, uint8
from typing import Optional

COMMON_QUANTIZATION_MATRIX = array(
    [
        [16, 11, 10, 16, 24, 40, 51, 61],
        [12, 12, 14, 19, 26, 58, 60, 55],
        [14, 13, 16, 24, 40, 57, 69, 56],
        [14, 17, 22, 29, 51, 87, 80, 62],
        [18, 22, 37, 56, 68, 109, 103, 77],
        [24, 35, 55, 64, 81, 104, 113, 92],
        [49, 64, 78, 87, 103, 121, 120, 101],
        [72, 92, 95, 98, 112, 100, 103, 99],
    ],
    dtype=uint8,
)

def quantize_forward(
    source: NDArray[float64],
    matrix: Optional[NDArray[uint8]] = None,
) -> NDArray[int64]:
    from numpy import atleast_2d, float64, int64, uint8

    source = atleast_2d(source).astype(float64)
    matrix = (
        atleast_2d(matrix).astype(uint8)
        if matrix
        else COMMON_QUANTIZATION_MATRIX
    )

    target = source.copy().astype(int64)
    for i in range(0, target.shape[0] - matrix.shape[0] + 1, matrix.shape[0]):
        for j in range(
            0, target.shape[1] - matrix.shape[1] + 1, matrix.shape[1]
```

```

        ):
            target[i : i + matrix.shape[0], j : j + matrix.shape[1]] = (
                source[i : i + matrix.shape[0], j : j + matrix.shape[1]]
                / matrix
            ).round()
    return target

def quantize_inverse(
    source: NDAarray[int64],
    matrix: Optional[NDAarray[uint8]] = None,
) -> NDAarray[float64]:
    from numpy import atleast_2d, float64, int64, uint8

    source = atleast_2d(source).astype(int64)
    matrix = (
        atleast_2d(matrix).astype(uint8)
        if matrix
        else COMMON_QUANTIZATION_MATRIX
    )

    target = source.copy().astype(float64)
    for i in range(0, target.shape[0] - matrix.shape[0] + 1, matrix.shape[0]):
        for j in range(
            0, target.shape[1] - matrix.shape[1] + 1, matrix.shape[1]
        ):
            target[i : i + matrix.shape[0], j : j + matrix.shape[1]] = (
                source[i : i + matrix.shape[0], j : j + matrix.shape[1]]
                * matrix
            )
    return target

```

Test the quantization and inverse quantization functions:

```

In [7]: from numpy.random import randn
        from numpy import set_printoptions

        mat = randn(17, 8) * 100
        mat_q = quantize_forward(mat)
        mat_iq = quantize_inverse(mat_q)

        set_printoptions(suppress=True)

        print(abs((mat - mat_iq)).mean())
        print(mat_q)

```

12.949407959340158

```
[[ 12  4 -17 -2  4 -1  3  1]
 [ -5 -2 -3 -1 -4 -2  0 -1]
 [ -3 -4  0  7  0  0  1  2]
 [  1 -4  2  5 -4  0  1  1]
 [  5  6  4  0 -2  1 -1  1]
 [ -1 -4  1  1  0  1  0  1]
 [  3  1 -2  0  1  1  1  1]
 [  1  0  0  0  1  0 -1  0]
 [  0  8 -9 13  0  1 -1 -1]
 [ -5 -11 10 -3  1 -2 -1  3]
 [  2  1 14 -2 -2 -2  1  2]
 [ 10  6 -1  0  1  0  0 -6]
 [  2  1  3  2  2  0  0  1]
 [ -1  0  1  3 -2  1  0  0]
 [  2  0 -1 -1  0  0  2  0]
 [  0 -1  2 -2  0  1 -1  1]
 [ 20 20 -24 183 67 193 -60 -19]]
```

Load the image and apply the DCT, quantization, inverse quantization, and IDCT functions to all the 8x8 blocks:

```
In [8]: from IPython.display import display, Markdown
from pathlib import Path
from PIL import Image
from numpy import asarray
from skimage.metrics import structural_similarity as ssim



source_path = Path("../resources/foreman_qcif_0_rgb.bmp").resolve()
source_image = Image.open(source_path).convert(mode="L")
source_path = Path("images/foreman_qcif_0_rgb.luma.source.bmp").resolve()
source_data = asarray(source_image)
transformed_data = dct_forward(source_data, 8)
quantized_data = quantize_forward(transformed_data)
dequantized_data = quantize_inverse(quantized_data)
target_data = dct_inverse(dequantized_data, 8)
target_image = Image.fromarray(target_data, mode="L")
target_path = Path("images/foreman_qcif_0_rgb.luma.target.2.bmp").resolve()
fidelity_ssim = ssim(source_data, target_data)

table_view = f"""\
| Variable Name | Value |
|-----|-----|
| source_image | ![source_image]({source_path}) |
| target_image | ![target_image]({target_path}) |
| fidelity_ssim | `{fidelity_ssim:.4f}` |
"""

for var_name in [
    "source_data",
    "transformed_data",
    "quantized_data",
    "dequantized_data",
    "target_data",
]:
    var_value = locals()[var_name]
    table_view += (
        f"| {var_name} | `{str(var_value).replace(\"\\n\", \"`<br>`\")}` |\\n"
    )

source_image.save(source_path)
```

```
target_image.save(target_path)
display(Markdown(table_view))
```

Variable Name	Value
source_image	
target_image	
fidelity_ssim	0.9418
source_data	[[32 233 ... 230 203] [39 212 ... 226 203] ... [15 132 ... 193 178] [14 131 ... 127 118]]
transformed_data	[[1531.5 -326.07 ... -16.01 1.28] [24.96 96.18 ... -6.66 7.32] ... [-1.68 1.61 ... 0.76 0.98] [2.79 -1.92 ... -5.4 -1.23]]
quantized_data	[[96 -30 ... 0 0] [2 8 ... 0 0] ... [0 0 ... 0 0] [0 0 ... 0 0]]
dequantized_data	[[1536. -330. ... 0. 0.] [24. 96. ... 0. 0.] ... [0. 0. ... 0. 0.] [0. 0. ... 0. 0.]]
target_data	[[22 242 ... 215 207] [43 203 ... 224 204] ... [13 132 ... 194 194] [13 132 ... 121 114]]